

Research Protocols & Standard Operating Procedures



"Slow is Smooth, Smooth is Fast."

This document defines the engineering standards required to contribute to this project. Violations of these protocols lead to "Code Rot" and are unacceptable.

1. The "Experiment-First" Architecture

We do not use a monolithic `src/` folder. We use an **Experiment-First** structure.

The Standard Layout

Every experiment MUST live in `experiments/v{VERSION}_{DESCRIPTION}/`.

Inside that folder, it MUST be self-contained:

- `models/`: A COPY of the model file used. Do not import from previous versions.
- `logs/`: Training logs.
- `audits/`: JSON results and validation reports.
- `tests/`: Unit tests and diagnostic scripts specific to this version.
- `journal.md`: The narrative log of the experiment (The "Captain's Log").
- `train.py`: The entry point script.

[!IMPORTANT]

The Journal is the Single Source of Truth

Everything about an experiment should be traceable from `journal.md`:

- What was tried (task breakdown)
- Why it was tried (hypothesis, methods)
- What happened (results, observations)
- What was learned (walkthroughs, proof of work)
- What's next (recommendations)

When you finish an experiment, merge `task.md` and `walkthrough.md` into journal sections. Future agents should only need to read `journal.md` to understand the complete story.

Why?

This allows us to delete or archive any old experiment without breaking the current SOTA. It decouples progress.

B. The Global Root Strategy

While `experiments/` is for *change*, the Root is for *permanence*.

- **scripts/**: V-Agnostic Utilities. Data preparation, global plotting, or repo maintenance tools. (If it imports a specific model, it belongs in `experiments/`).

- **docs/**: **Deep Knowledge**. Mathematical proofs, architectural whitepapers, and guides that outlive any single version.
- **data/**: **Immutable Source**. Raw datasets. Experiments treat this as Read-Only.
- **configs/**: **Shared Baselines**. Global configuration constants (if needed).
- **final_report/**: **The Publication**. When a milestone is reached, polished figures and summaries are copied here for external consumption (e.g. arXiv/GitHub pages).

C. Strategic Continuity (The Roadmap)

While `journal.md` is for the **Experiment**, `ROADMAP.md` is for the **Project**.

- **Requirement**: Every project MUST have a `ROADMAP.md` in the root (initialized from `ROADMAP TEMPLATE.md`).
- **The Protocol**:
 - Planning**: Before starting a new version, check the active Phase in the roadmap.
 - Execution**: If an experiment triggers a strategic pivot, update the roadmap immediately.
 - Handoff**: Before closing a session, update the Roadmap's Phase table with your findings.

Why? To prevent "Short-Sighted Agent" syndrome. The Roadmap ensures results aren't just logged, but are integrated into the project's long-term trajectory.

2. The Agentic Workflow

A. Matrix Mode (Deep Work)

When entering a complex task:

1. **Plan**: Create `implementation_plan.md`.
2. **Execute**: Write code, strictly following the plan.
3. **Verify**: Run `audit.py` or `tests/`.
4. **Document**: Update `journal.md` and `walkthrough.md`.

B. The "Mini-Train" Rule

Before launching a full run:

1. Run the training script for **200 steps** (or 1 epoch on small data).
2. **Verify Convergence**: Does loss go down?
3. **Verify Artifacts**: Are checkpoints saving? Are logs writing?
Never launch a 2-day job without a 2-minute test.

C. The "Mini-Audit" Rule

Before claiming success or running a benchmark:

1. **Constraint**: Must run in **< 5 minutes**.
2. **Task**: A simplified, synthetic version of the Hard Task (e.g. "Toy Recall" vs "Book Recall").
3. **Threshold**: Must hit **100% Accuracy** (or equivalent 0.0 Loss).

If it fails the Toy Task, it will fail the Real Task. Fail fast.

3. Code Hygiene

- **Explicit Imports:** Avoid `from model import *`.
- **Config Separation:** Hyperparameters should be at the top of the script or in a config object, not buried in loops.
- **Assertion Guardrails:** Use `assert` statements aggressively to validate tensor shapes at the start of `forward()`.

4. Final Report Protocol

When an experiment concludes:

1. **Success:** Mark as SOTA in `LEADERBOARD.md`. Merge findings into `README.md`.
 2. **Failure:** Document *why* in `journal.md`. Move folder to `archive/` if it is clutter.
 3. **Handoff:** Write a summary in `journal.md` for the next agent.
-

5. Git Hygiene & Version Control

A. The "Heavy File" Ban

Never commit:

- Checkpoints (`.pt`, `.ckpt`)
 - Large Logs (`.log`, `.txt` > 1MB)
 - Datasets (`.bin`, `.jsonl`)
- Why?**: Bloats the repo size forever. Use `.gitignore` religiously.

B. The "folder-as-branch" Strategy

In this research repo, we prefer **New Folders** (`experiments/vXX`) over **Git Branches** for experimental variations.

- **Benefit:** You can diff `v1/models/model.py` vs `v2/models/model.py` directly in the editor.
- **Benefit:** Multiple experiments can run simultaneously on the same machine without switching git branches (which would break running jobs).

C. The "Clean Start" Rule

Before launching a training run:

1. **Commit your code.**
2. Log the **Git Hash** in your `journal.md` or training log.

This ensures that if the run is a success 3 days later, you know EXACTLY what code produced it.

6. Agent Cleanup Protocol

When: After completing a discrete task (e.g., reorganization, audit, migration).

A. Definition: "Stub Files"

Files created **only** to perform a one-time operation:

- `cleanup_*.py` (e.g., `cleanup_experiments.py`)
- `audit_structure.py` (after generating the report)
- `migrate_*.py`
- `temp_*.py`

NOT stub files:

- `train.py, audit.py` (persistent entry points)
- Scripts in `experiments/vXX/tests/` (diagnostic tools)
- Scripts in `scripts/` (reusable utilities)

B. The Rule

Before calling `notify_user` to complete a task:

1. Identify stub files created during this session.
2. Delete them: `rm cleanup_experiments.py audit_structure.py`
3. Commit the clean state.

C. Rationale

- **Avoid Clutter:** The root should contain **only** persistent infrastructure.
- **Prevent Confusion:** Future agents won't mistakenly re-run stale cleanup scripts.
- **Signal Completion:** A clean directory signals "this task is done."

D. Artifact Migration

At the end of an experiment, migrate planning/verification artifacts from `/artifacts/` to the experiment folder:

Step 1: Move Artifacts

```
# If you created task.md or walkthrough.md in artifacts directory  
mv /path/to/artifacts/task.md experiments/v68/task.md  
mv /path/to/artifacts/walkthrough.md experiments/v68/walkthrough.md
```

Step 2: Merge into journal.md (Recommended)

Instead of keeping separate files, merge content into `journal.md` sections:

```
# experiments/v68/journal.md  
  
## Task Breakdown  
[Paste content from task.md here]  
  
## Walkthrough  
[Paste content from walkthrough.md here - proof of work, validation]
```

Then delete standalone files: `rm experiments/v68/task.md experiments/v68/walkthrough.md`

Alternative: Keep Separate + Symlink (If you need them accessible from root)

```
# Keep originals in experiment folder
# Create symlinks in artifacts for convenience
ln -s $(pwd)/experiments/v68/task.md /path/to/artifacts/task_v68.md
```

Why merge? Single source of truth, no sync issues, complete narrative in one file.

7. Research Methodology Protocols

A. The Smoke Test Mandate

When: Creating any script with **moderate or higher dependency complexity**.

Complexity Indicators:

- Inherits from multiple classes
- Heavily parametrized (> 5 config variables)
- Interacts with GPU/CUDA
- Loads external models or data

The Rule: Create `tests/smoke_test_[FEATURE].py` alongside the main script.

Example:

```
# tests/smoke_test_v68_lane_unification.py
import torch
from models.hybrid_transformer_v68 import HybridTransformerV68

def test_initialization():
    model = HybridTransformerV68(vocab_size=100, hidden_size=64)
    assert model is not None

def test_forward_pass():
    model = HybridTransformerV68(vocab_size=100, hidden_size=64)
    x = torch.randint(0, 100, (2, 10)) # Batch=2, Seq=10
    logits, _ = model(x)
    assert logits.shape == (2, 10, 100)

if __name__ == "__main__":
    test_initialization()
    test_forward_pass()
    print("✅ All smoke tests passed.")
```

Why: Catches import errors, shape bugs, and device mismatches **before** launching a 2-day training run.

B. The Metrics Constitution

Every project must define its **North Star Metrics** early and evaluate them **consistently**.

Step 1: Define

Create `METRICS_METHODOLOGY.md` specifying:

- **Core Metrics** (e.g., Recall, Gate Density, Throughput)
- **Success Thresholds** (e.g., Recall > 90%, Density < 15%)
- **Evaluation Protocol** (e.g., batch size, context length, hardware)

Step 2: Enforce via Class Inheritance

Build a base audit class that all experiments inherit:

```
# experiments/base_audit.py
class BaseAudit:
    def __init__(self, model, dataset):
        self.model = model
        self.dataset = dataset

    def compute_recall(self):
        raise NotImplementedError

    def compute_density(self):
        raise NotImplementedError

    def run_full_audit(self):
        results = {
            "recall": self.compute_recall(),
            "density": self.compute_density(),
            # ... other metrics
        }
        return results

# experiments/v68_lane_unification/audit_v68.py
from experiments.base_audit import BaseAudit

class AuditV68(BaseAudit):
    def compute_recall(self):
        # V68-specific recall logic
        pass
```

Why: Ensures every experiment reports the **same metrics** in the **same format**, enabling apples-to-apples comparison in `LEADERBOARD.md`.

C. The Principled Problem-Solving Framework

When you hit a major architectural or algorithmic impasse:

1. Establish Canonical Wisdom

- **Survey the Literature:** Find 3-5 foundational papers addressing similar problems.
- **Document Best Practices:** In a `docs/[TOPIC].md` file (e.g., `docs/gradient_flow.md`, `docs/memory_gating.md`).
- **Extract Heuristics:** E.g., "Weight Update Ratio should be ~ 10^{-3} " (from DEBUGGING_GUIDE.md).

2. Ground Your Approach

- **Cite Inspiration:** "Our regret gating is inspired by Hindsight Experience Replay (Andrychowicz et al., 2017)."
- **Reference Implementations:** "We adapted the attention masking from Transformer-XL's official repo."
- **Benchmark Against Empirical Results:** "Paper X reports 85% recall at 100k context. Our baseline must exceed this."

3. Document in the Journal

Every experiment's `journal.md` **must** include a **Methods** section:

```
## Methods

### Architectural Decisions
1. **Lane Unification**: Inspired by [Paper X], we merge Lexical/Semantic lanes using Centroid-based addressing.
2. **Gating Logic**: Following [Paper Y]'s finding that surprisal-based gates reduce redundancy, we implement...

### Key References
- [Andrychowicz et al., 2017] Hindsight Experience Replay
- [Wu et al., 2022] Memorizing Transformers
- [Bulatov et al., 2023] Recurrent Memory Transformer

### Deviations from Literature
We use Z-score normalization instead of raw entropy because [REASON].
```

Why:

- **Reproducibility:** Future agents (or humans) can understand your reasoning.
- **Credibility:** Reviewers see you're not reinventing the wheel blindly.
- **Velocity:** You avoid dead-ends that prior work already explored.

D. The First-Stop-Shop Rule

Before implementing a new feature:

1. **Check the paper's official repo** (if available).
2. **Run their test suite** to understand expected behavior.
3. **Copy their hyperparameters** as your baseline.
4. **Only diverge** if you have a hypothesis for why your approach is better.

Anti-Pattern: "I'll just implement it from scratch based on the math in the paper."

Result: 3 days of debugging shape errors that the official repo already solved.

Correct Pattern: "I cloned their repo, confirmed it works, adapted their `forward()` to my use case."

Result: Working implementation in 2 hours.

7. Multi-Environment & Hardware Layout

We operate in a **Hybrid Compute** environment.

A. The Agent Environment (GCP Compute Engine / L4)

- **Role: Home Base.** Agents live here.
- **Capabilities:** Coding, Logic, Mini-Trains (L4 GPU, **22GB vRAM**).
- **Rule:** ALL code is written and versioned here.

B. The external "Heavy Lift" (Google Colab / A100 Cluster)

- **Role: Muscle.** User manually runs massive jobs here.
- **Capabilities:** High-RAM Training (A100 GPU, **80GB vRAM**).
- **Process:**
 1. **Export:** User zips the code from Agent Env.
 2. **Train:** User runs job on Colab.
 3. **Import:** User uploads logs/checkpoints back to Agent Env via `gdown`.

C. The "Airlock" Protocol

When handling external results:

1. **Never** manually edit code on Colab and forget to update the Agent.
2. **Always** bring the `log.txt` back to `experiments/vXX/logs/` so the Agent can analyze it.