

```
eggfm/config.py      Sat Dec 06 23:26:45 2025      1
# src/mantra/eggfm/config.py

from __future__ import annotations

from mantra.config import EnergyModelConfig, EnergyTrainConfig, EnergyModelBundle

__all__ = [
    "EnergyModelConfig",
    "EnergyTrainConfig",
    "EnergyModelBundle",
]

```

```
# AnnDataPyTorch.py

import numpy as np
import torch
from torch.utils.data import Dataset
from scipy import sparse
import scanpy as sc # only for type hints; not strictly necessary

class AnnDataExpressionDataset(Dataset):
    """
    Wraps an AnnData object's X matrix (after prep()) as a PyTorch dataset.
    Uses HVG, log-normalized expression directly.
    """

    def __init__(self, X, float_dtype=np.float32):
        if sparse.issparse(X):
            X = X.toarray()
        X = np.asarray(X, dtype=float_dtype)

        mean = X.mean(axis=0, keepdims=True)
        std = X.std(axis=0, keepdims=True)

        # prevent divide-by-zero or tiny variance explosions
        std = np.clip(std, 1e-2, None)

        # store for later (without the extra batch dim)
        self.mean = mean.astype(float_dtype).squeeze(0) # shape [D]
        self.std = std.astype(float_dtype).squeeze(0) # shape [D]

        self.X = (X - mean) / std

    def __len__(self) -> int:
        return self.X.shape[0]

    def __getitem__(self, idx: int) -> torch.Tensor:
        return torch.from_numpy(self.X[idx])
```

```
# src/mantra/eggfm/inference.py

from __future__ import annotations

from pathlib import Path
from typing import Any, Dict, List, Optional, Sequence, Union

import numpy as np
import torch
from torch import nn, Tensor

from mantra.eggfm.models import EnergyMLP


class EnergyScorer:
    """
    Wraps a trained EnergyMLP + normalization (+ optional projection)
    so we can compute energies in a consistent way.

    Supports:
    - HVG / gene space: x -> normalize -> E(x)
    - Embedding space: x -> project (PCA) -> normalize -> E_z(z)
    """

    def __init__(
        self,
        energy_model: nn.Module,
        mean: Optional[Tensor],
        std: Optional[Tensor],
        var_names: Optional[Sequence[str]] = None,
        proj_matrix: Optional[Tensor] = None,      # [G, d] for embedding case
        space: str = "hvg",
        device: Optional[torch.device] = None,
    ) -> None:
        self.device = device or torch.device(
            "cuda" if torch.cuda.is_available() else "cpu"
        )

        self.energy_model = energy_model.to(self.device)
        self.energy_model.eval()
        for p in self.energy_model.parameters():
            p.requires_grad_(False)

        self.space = space

        # mean/std are always in the *model feature space*: [D_model]
        self.mean = None if mean is None else mean.to(self.device).view(1, -1)
        self.std = None if std is None else std.to(self.device).view(1, -1)

        # gene feature metadata (for HVG space alignment; optional)
        self.var_names: Optional[List[str]] = None
        self._name_to_idx: Optional[Dict[str, int]] = None
        if var_names is not None:
            self.var_names = [str(v) for v in var_names]
            self._name_to_idx = {name: i for i, name in enumerate(self.var_names)}

        # optional projection (for embedding case): [G_raw, D_model]
        self.proj_matrix: Optional[Tensor] = None
        if proj_matrix is not None:
            proj_matrix = proj_matrix.to(self.device)
            self.proj_matrix = proj_matrix

    # -----
    # Construction helper
    # -----
```

```
cls,
ckpt_path: Union[str, Path],
device: Optional[torch.device] = None,
) -> "EnergyScorer":
"""
Load an EnergyScorer from a .pt checkpoint.

Expects ckpt to contain something like:

{
    "state_dict": ...,
    "model_cfg": {"hidden_dims": [...]},
    "n_genes": int,           # D_model
    "space": "hvg" or "embedding",
    "var_names": [...],      # optional, for HVG space alignment
    "mean": ...,
    "std": ...,
    # optional for embedding:
    "proj_matrix": np.ndarray [G_raw, D_model],
}
"""

ckpt_path = Path(ckpt_path)
ckpt = torch.load(ckpt_path, map_location=device or "cpu")

n_genes = ckpt.get("n_genes")
model_cfg = ckpt.get("model_cfg", {})
space = ckpt.get("space", "hvg")

# reconstruct EnergyMLP in model feature space
energy_model = EnergyMLP(
    n_genes=n_genes,
    **model_cfg,
)
energy_model.load_state_dict(ckpt["state_dict"])

def _to_tensor_or_none(key: str) -> Optional[Tensor]:
    if key not in ckpt or ckpt[key] is None:
        return None
    arr = ckpt[key]
    if isinstance(arr, Tensor):
        return arr
    return torch.as_tensor(arr, dtype=torch.float32)

mean = _to_tensor_or_none("mean")
std = _to_tensor_or_none("std")
var_names = ckpt.get("var_names", None)

proj_matrix = _to_tensor_or_none("proj_matrix")  # for embedding space

return cls(
    energy_model=energy_model,
    mean=mean,
    std=std,
    var_names=var_names,
    proj_matrix=proj_matrix,
    space=space,
    device=device,
)

# -----
# Internal helpers
# -----


def _ensure_tensor(self, x: Union[Tensor, np.ndarray]) -> Tensor:
    if isinstance(x, Tensor):
        return x.to(self.device, dtype=torch.float32)
    else:
        return torch.as_tensor(x, dtype=torch.float32, device=self.device)
```

```
def _reorder_by_genes(
    self,
    x: Tensor, # [B, G_in]
    gene_names: Optional[Sequence[str]],
) -> Tensor:
    """
    If var_names and gene_names are provided, reorder x to match training order.
    Otherwise, assume x is already aligned.
    """
    if self.var_names is None or gene_names is None:
        return x

    if len(self.var_names) != x.shape[1]:
        raise ValueError(
            f"EnergyScorer: mismatch between model gene dim ({len(self.var_names)}) "
            f"and input x.shape[1] ({x.shape[1]})."
        )

    input_name_to_idx = {str(name): i for i, name in enumerate(gene_names)}

    try:
        indices = [input_name_to_idx[name] for name in self.var_names]
    except KeyError as e:
        missing = str(e.args[0])
        raise KeyError(
            f"EnergyScorer: gene {missing} from training var_names "
            f"not found in provided gene_names."
        )

    idx = torch.as_tensor(indices, dtype=torch.long, device=x.device)
    return x[:, idx]

def _apply_normalization(self, z: Tensor) -> Tensor:
    if self.mean is None or self.std is None:
        return z
    return (z - self.mean) / self.std

# -----
# Public API
# -----
```

```
@torch.no_grad()
def score(
    self,
    x_raw: Union[Tensor, np.ndarray],
    gene_names: Optional[Sequence[str]] = None,
) -> Tensor:
    """
    Compute energy for a batch of states x_raw in *gene space*:
    - If space == "hvg": x_raw is already in model feature space (after alignment).
    - If space == "embedding": x_raw is in gene space; we project with proj_matrix.

    Returns energies: [B].
    """
    x = self._ensure_tensor(x_raw) # [B, G_in]
    x = self._reorder_by_genes(x, gene_names) # optionally align to var_names

    # If we have a projection matrix, go to embedding space
    if self.proj_matrix is not None:
        z = x @ self.proj_matrix # [B, D_model]
    else:
        z = x # [B, D_model]

    z_norm = self._apply_normalization(z) # [B, D_model]
    energy = self.energy_model(z_norm) # [B] or [B, 1]
    if energy.ndim == 2:
```

```
        energy = energy.squeeze(-1)
    return energy

@torch.no_grad()
def score_delta(
    self,
    x_ref: Union[Tensor, np.ndarray],
    deltaE_pred: Union[Tensor, np.ndarray],
    gene_names: Optional[Sequence[str]] = None,
) -> Tensor:
    """
    Convenience: score energy of x_hat = x_ref + \224E_pred.

    x_ref: [G] or [1,G]
    deltaE_pred: [B,G]
    """
    x_ref_t = self._ensure_tensor(x_ref)
    if x_ref_t.ndim == 1:
        x_ref_t = x_ref_t.unsqueeze(0) # [1,G]
    delta_t = self._ensure_tensor(deltaE_pred) # [B,G]

    if x_ref_t.shape[1] != delta_t.shape[1]:
        raise ValueError(
            f"x_ref dim {x_ref_t.shape[1]} != deltaE_pred dim {delta_t.shape[1]}"
        )

    x_hat = x_ref_t + delta_t
    return self.score(x_hat, gene_names=gene_names)
```

```
# EnergyMLP.py

from typing import Sequence, Optional
import torch
from torch import nn

class EnergyMLP(nn.Module):
    """
    E(x) = <E_theta(x), x> where E_theta is an MLP with nonlinearities.

    x is HVG, log-normalized expression (optionally mean-centered).

    We also expose a latent representation z(x) from the last hidden layer,
    which can be used as a geometry for manifold learning.
    """

    def __init__(self,
                 n_genes: int,
                 hidden_dims: Sequence[int] = (512, 512, 512, 512),
                 activation: Optional[nn.Module] = None,
                 ) :
        super().__init__()
        if activation is None:
            activation = nn.Softplus()

        layers = []
        in_dim = n_genes
        for h in hidden_dims:
            layers.append(nn.Linear(in_dim, h))
            layers.append(activation)
            in_dim = h

        # encoder: maps x \in R^{n_genes} to z \in R^{hidden_dims[-1]}
        self.hidden = nn.Sequential(*layers)

        # head: maps z \in R^{hidden_dims[-1]} to v(z) \in R^{n_genes}
        self.vector_head = nn.Linear(in_dim, n_genes)

        # store for convenience
        self.n_genes = n_genes
        self.latent_dim = in_dim

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """
        Standard forward used in training:
        x: (B, D)
        returns: energy (B, )
        """
        if x.dim() == 1:
            x = x.unsqueeze(0)

        z = self.hidden(x)                      # (B, latent_dim)
        v = self.vector_head(z)                  # (B, D)
        energy = (v * x).sum(dim=-1)           # <v(x), x>
        return energy

    @torch.no_grad()
    def encode(self, x: torch.Tensor) -> torch.Tensor:
        """
        Return latent representation z(x) from the last hidden layer.
        x: (B, D)
        returns: z (B, latent_dim)
        """
        if x.dim() == 1:
            x = x.unsqueeze(0)
        z = self.hidden(x)
```

```
    return z

def score(self, x: torch.Tensor) -> torch.Tensor:
    """
    score(x) = \log p(x) = -E(x)
    """
    x = x.clone().detach().requires_grad_(True)
    energy = self.forward(x) # (B, )
    energy_sum = energy.sum()
    (grad,) = torch.autograd.grad(
        energy_sum,
        x,
        create_graph=False,
        retain_graph=False,
        only_inputs=True,
    )
    score = -grad
    return score
```

```

eggfm/trainer.py      Sun Dec 07 02:21:38 2025      1
# src/mantra/eggfm/trainer.py

from __future__ import annotations

from typing import Optional

import numpy as np
import torch
import scanpy as sc
from torch import optim
from torch.utils.data import DataLoader

from mantra.eggfm.models import EnergyMLP
from mantra.eggfm.dataset import AnnDataExpressionDataset
from mantra.config import EnergyModelConfig, EnergyTrainConfig, EnergyModelBundle


class EnergyTrainer:
    """
    Denoising score-matching trainer for EnergyMLP.

    Given:
    - model
    - standardized dataset (AnnDataExpressionDataset)
    - EnergyTrainConfig

    It runs the DSM loop and returns the best-trained model.
    """

    def __init__(
        self,
        model: EnergyMLP,
        dataset: AnnDataExpressionDataset,
        train_cfg: EnergyTrainConfig,
    ) -> None:
        self.model = model
        self.dataset = dataset
        self.train_cfg = train_cfg

        device_str = train_cfg.device or ("cuda" if torch.cuda.is_available() else "cpu")
        self.device = torch.device(device_str)
        self.model.to(self.device)

        self.loader = DataLoader(
            dataset,
            batch_size=train_cfg.batch_size,
            shuffle=True,
            drop_last=True,
        )

        self.optimizer = optim.Adam(
            self.model.parameters(),
            lr=float(train_cfg.lr),           # force-cast in case YAML gave a string
            weight_decay=float(train_cfg.weight_decay),
        )

        self.best_loss: float = float("inf")
        self.best_state_dict: Optional[dict] = None

    def train(self) -> EnergyMLP:
        sigma = float(self.train_cfg.sigma)
        grad_clip = float(self.train_cfg.grad_clip)
        early_stop_patience = int(self.train_cfg.early_stop_patience)
        early_stop_min_delta = float(self.train_cfg.early_stop_min_delta)
        num_epochs = int(self.train_cfg.num_epochs)

        self.model.train()
        epochs_without_improve = 0

```

```
n_total = len(self.dataset)

for epoch in range(num_epochs):
    running_loss = 0.0

    for xb in self.loader:
        xb = xb.to(self.device) # (B, D), already standardized

        # Sample Gaussian noise
        eps = torch.randn_like(xb)
        y = xb + sigma * eps
        y.requires_grad_(True)

        # Energy and score
        energy = self.model(y)           # (B, )
        energy_sum = energy.sum()       # scalar

        (grad_y,) = torch.autograd.grad(
            energy_sum,
            y,
            create_graph=True,
            retain_graph=True,
            only_inputs=True,
        )
        s_theta = -grad_y

        # DSM target: -(y - x) / sigma^2
        target = -(y - xb) / (sigma**2)

        # MSE over batch and dimensions
        loss = ((s_theta - target) ** 2).sum(dim=1).mean()

        self.optimizer.zero_grad()
        loss.backward()
        if grad_clip > 0.0:
            torch.nn.utils.clip_grad_norm_(
                self.model.parameters(),
                grad_clip,
            )
        self.optimizer.step()

        running_loss += loss.item() * xb.size(0)
    epoch_loss = running_loss / n_total
    print(
        f"[Energy DSM] Epoch {epoch+1}/{num_epochs} loss={epoch_loss:.6e}",
        flush=True,
    )

    improved = epoch_loss + early_stop_min_delta < self.best_loss
    if improved:
        self.best_loss = epoch_loss
        self.best_state_dict = self.model.state_dict()
        epochs_without_improve = 0
    else:
        epochs_without_improve += 1

    if early_stop_patience > 0 and epochs_without_improve >= early_stop_patience:
        print(
            f"[Energy DSM] Early stopping at epoch {epoch+1} "
            f"(best_loss={self.best_loss:.6e})",
            flush=True,
        )
        break

if self.best_state_dict is not None:
    self.model.load_state_dict(self.best_state_dict)
```

```
        return self.model

# -----
# High-level convenience wrapper: AnnData -> EnergyModelBundle
# -----


def train_energy_model(
    ad_prep: sc.AnnData,
    model_cfg: EnergyModelConfig,
    train_cfg: EnergyTrainConfig,
) -> EnergyModelBundle:
    """
    Convenience wrapper used by scripts:

    AnnData -> AnnDataExpressionDataset -> EnergyMLP -> EnergyTrainer

    Trains an energy-based model on preprocessed AnnData using DSM
    and returns an EnergyModelBundle.
    """
    # ----- dataset: HVG or PCA -----
    latent_space = "hvg" # promote to config later if you want
    if latent_space == "hvg":
        X = ad_prep.X
    else:
        if "X_pca" not in ad_prep.obsm:
            sc.pp.pca(ad_prep, n_comps=50)
        X = ad_prep.obsm["X_pca"]

    dataset = AnnDataExpressionDataset(X)
    n_genes = dataset.X.shape[1]

    # record normalization
    mean = dataset.mean # [D]
    std = dataset.std # [D]
    feature_names = np.array(ad_prep.var_names)

    # ----- model -----
    hidden_dims = tuple(model_cfg.hidden_dims)
    model = EnergyMLP(
        n_genes=n_genes,
        hidden_dims=hidden_dims,
    )

    # ----- trainer -----
    trainer = EnergyTrainer(
        model=model,
        dataset=dataset,
        train_cfg=train_cfg,
    )
    best_model = trainer.train()

    return EnergyModelBundle(
        model=best_model,
        mean=mean,
        std=std,
        feature_names=feature_names,
        space=latent_space,
    )
)
```

```
eggfm/__init__.py           Sun Dec 07 02:07:08 2025      1
# src/mantra/eggfm/__init__.py

from .models import EnergyMLP
from .dataset import AnnDataExpressionDataset
from .trainer import EnergyTrainer, train_energy_model
from .inference import EnergyScorer

__all__ = [
    "EnergyMLP",
    "AnnDataExpressionDataset",
    "EnergyTrainer",
    "train_energy_model",
    "EnergyScorer",
]

```