

```
#!/usr/bin/env python
# src/mantra/eggfm/config.py
"""

Configuration containers for EGGFM (Energy-Guided Geometric Flow Model).

EnergyModelConfig:
    - architecture hyperparameters for the energy network (EnergyMLP)
EnergyTrainConfig:
    - DSM training hyperparameters and subset options
EnergyModelBundle:
    - trained model + normalization metadata for downstream use
"""

from __future__ import annotations

from dataclasses import dataclass
from typing import Sequence, Optional, List

import numpy as np
from torch import nn

@dataclass
class EnergyModelConfig:
    """
    Architecture config for the energy network (e.g. EnergyMLP).

    This is deliberately small and explicit so that the model can be fully
    reconstructed from a YAML block or CLI + defaults.
    """

    # Hidden layer widths for the MLP (input/output inferred from data)
    hidden_dims: Sequence[int] = (512, 512, 512, 512)

@dataclass
class EnergyTrainConfig:
    """
    Hyperparameters for denoising score matching (DSM) training of EGGFM.

    Controls batch size, LR, DSM noise scale, regularization, and optional
    subsampling of cells / HVGs.
    """

    # Core training loop
    batch_size: int = 2048
    num_epochs: int = 50
    lr: float = 1e-4

    # DSM noise scale (Gaussian corruption std)
    sigma: float = 0.1

    # Regularization / stabilization
    weight_decay: float = 0.0
    grad_clip: float = 0.0 # 0.0 = disabled

    # Early stopping on DSM loss (0 patience = disabled)
    early_stop_patience: int = 0
    early_stop_min_delta: float = 0.0

    # Device + subsampling
    device: Optional[str] = None          # e.g. "cuda", "cpu", or None -> auto
    n_cells_sample: Optional[int] = None   # if set, sample this many cells per epoch
    max_hvg: Optional[int] = None          # if set, restrict to top-N HVGs for EGGFM

@dataclass
class EnergyModelBundle:
```



```
# AnnDataPyTorch.py

import numpy as np
import torch
from torch.utils.data import Dataset
from scipy import sparse

class AnnDataExpressionDataset(Dataset):
    """
    Wraps an AnnData object's X matrix (after prep()) as a PyTorch dataset.
    Uses HVG, log-normalized expression directly.
    """

    def __init__(self, X, float_dtype=np.float32):
        if sparse.issparse(X):
            X = X.toarray()
        X = np.asarray(X, dtype=float_dtype)

        mean = X.mean(axis=0, keepdims=True)
        std = X.std(axis=0, keepdims=True)

        # prevent divide-by-zero or tiny variance explosions
        std = np.clip(std, 1e-2, None)

        # store for later (without the extra batch dim)
        self.mean = mean.astype(float_dtype).squeeze(0)      # shape [D]
        self.std = std.astype(float_dtype).squeeze(0)         # shape [D]

        self.X = (X - mean) / std

    def __len__(self) -> int:
        return self.X.shape[0]

    def __getitem__(self, idx: int) -> torch.Tensor:
        return torch.from_numpy(self.X[idx])
```

```
# src/mantra/eggfm/inference.py

from __future__ import annotations

from pathlib import Path
from typing import Dict, List, Optional, Sequence, Union

import numpy as np
import torch
from torch import nn, Tensor

from mantra.eggfm.models import EnergyMLP


class EnergyScorer:
    """
    Wraps a trained EnergyMLP + normalization (+ optional projection)
    so we can compute energies in a consistent way.

    Supports:
    - HVG / gene space: x -> normalize -> E(x)
    - Embedding space: x -> project (PCA) -> normalize -> E_z(z)
    """

    def __init__(
        self,
        energy_model: nn.Module,
        mean: Optional[Tensor],
        std: Optional[Tensor],
        var_names: Optional[Sequence[str]] = None,
        proj_matrix: Optional[Tensor] = None,      # [G, d] for embedding case
        space: str = "hvg",
        device: Optional[torch.device] = None,
    ) -> None:
        self.device = device or torch.device(
            "cuda" if torch.cuda.is_available() else "cpu"
        )

        self.energy_model = energy_model.to(self.device)
        self.energy_model.eval()
        for p in self.energy_model.parameters():
            p.requires_grad_(False)

        self.space = space

        # mean/std are always in the *model feature space*: [D_model]
        self.mean = None if mean is None else mean.to(self.device).view(1, -1)
        self.std = None if std is None else std.to(self.device).view(1, -1)

        # gene feature metadata (for HVG space alignment; optional)
        self.var_names: Optional[List[str]] = None
        self._name_to_idx: Optional[Dict[str, int]] = None
        if var_names is not None:
            self.var_names = [str(v) for v in var_names]
            self._name_to_idx = {name: i for i, name in enumerate(self.var_names)}

        # optional projection (for embedding case): [G_raw, D_model]
        self.proj_matrix: Optional[Tensor] = None
        if proj_matrix is not None:
            proj_matrix = proj_matrix.to(self.device)
            self.proj_matrix = proj_matrix

    # -----
    # Construction helper
    # -----
```

```
cls,
ckpt_path: Union[str, Path],
device: Optional[torch.device] = None,
) -> "EnergyScorer":
"""
Load an EnergyScorer from a .pt checkpoint.

Expects ckpt to contain something like:

{
    "state_dict": ...,
    "model_cfg": {"hidden_dims": [...]},
    "n_genes": int,           # D_model
    "space": "hvg" or "embedding",
    "var_names": [...],      # optional, for HVG space alignment
    "mean": ...,
    "std": ...,
    # optional for embedding:
    "proj_matrix": np.ndarray [G_raw, D_model],
}
"""

ckpt_path = Path(ckpt_path)
ckpt = torch.load(ckpt_path, map_location=device or "cpu")

n_genes = ckpt.get("n_genes")
model_cfg = ckpt.get("model_cfg", {})
space = ckpt.get("space", "hvg")

# reconstruct EnergyMLP in model feature space
energy_model = EnergyMLP(
    n_genes=n_genes,
    **model_cfg,
)
energy_model.load_state_dict(ckpt["state_dict"])

def _to_tensor_or_none(key: str) -> Optional[Tensor]:
    if key not in ckpt or ckpt[key] is None:
        return None
    arr = ckpt[key]
    if isinstance(arr, Tensor):
        return arr
    return torch.as_tensor(arr, dtype=torch.float32)

mean = _to_tensor_or_none("mean")
std = _to_tensor_or_none("std")
var_names = ckpt.get("var_names", None)

proj_matrix = _to_tensor_or_none("proj_matrix")  # for embedding space

return cls(
    energy_model=energy_model,
    mean=mean,
    std=std,
    var_names=var_names,
    proj_matrix=proj_matrix,
    space=space,
    device=device,
)

# -----
# Internal helpers
# -----


def _ensure_tensor(self, x: Union[Tensor, np.ndarray]) -> Tensor:
    if isinstance(x, Tensor):
        return x.to(self.device, dtype=torch.float32)
    else:
        return torch.as_tensor(x, dtype=torch.float32, device=self.device)
```

```

def _reorder_by_genes(
    self,
    x: Tensor,                      # [B, G_in]
    gene_names: Optional[Sequence[str]],
) -> Tensor:
    """
    If var_names and gene_names are provided, reorder x to match training order.
    Otherwise, assume x is already aligned.
    """
    if self.var_names is None or gene_names is None:
        return x

    if len(self.var_names) != x.shape[1]:
        raise ValueError(
            f"EnergyScorer: mismatch between model gene dim ({len(self.var_names)}) "
            f"and input x.shape[1] ({x.shape[1]})."
        )

    input_name_to_idx = {str(name): i for i, name in enumerate(gene_names)}

    try:
        indices = [input_name_to_idx[name] for name in self.var_names]
    except KeyError as e:
        missing = str(e.args[0])
        raise KeyError(
            f"EnergyScorer: gene {missing} from training var_names "
            f"not found in provided gene_names."
        )

    idx = torch.as_tensor(indices, dtype=torch.long, device=x.device)
    return x[:, idx]

def _apply_normalization(self, z: Tensor) -> Tensor:
    if self.mean is None or self.std is None:
        return z
    return (z - self.mean) / self.std

# -----
# Public API
# -----


@torch.no_grad()
def score(
    self,
    x_raw: Union[Tensor, np.ndarray],
    gene_names: Optional[Sequence[str]] = None,
) -> Tensor:
    """
    Compute energy for a batch of states x_raw in *gene space*:

    - If space == "hvg": x_raw is already in model feature space (after alignment).
    - If space == "embedding": x_raw is in gene space; we project with proj_matrix.

    Returns energies: [B].
    """
    x = self._ensure_tensor(x_raw) # [B, G_in]
    x = self._reorder_by_genes(x, gene_names) # optionally align to var_names

    # If we have a projection matrix, go to embedding space
    if self.proj_matrix is not None:
        z = x @ self.proj_matrix # [B, D_model]
    else:
        z = x # [B, D_model]

    z_norm = self._apply_normalization(z) # [B, D_model]
    energy = self.energy_model(z_norm) # [B] or [B, 1]
    if energy.ndim == 2:

```

```
        energy = energy.squeeze(-1)
    return energy

@torch.no_grad()
def score_delta(
    self,
    x_ref: Union[Tensor, np.ndarray],
    deltaE_pred: Union[Tensor, np.ndarray],
    gene_names: Optional[Sequence[str]] = None,
) -> Tensor:
    """
    Convenience: score energy of x_hat = x_ref + \224E_pred.

    x_ref: [G] or [1,G]
    deltaE_pred: [B,G]
    """
    x_ref_t = self._ensure_tensor(x_ref)
    if x_ref_t.ndim == 1:
        x_ref_t = x_ref_t.unsqueeze(0) # [1,G]
    delta_t = self._ensure_tensor(deltaE_pred) # [B,G]

    if x_ref_t.shape[1] != delta_t.shape[1]:
        raise ValueError(
            f"x_ref dim {x_ref_t.shape[1]} != deltaE_pred dim {delta_t.shape[1]}"
        )

    x_hat = x_ref_t + delta_t
    return self.score(x_hat, gene_names=gene_names)
```

```
# EnergyMLP.py

from typing import Sequence, Optional
import torch
from torch import nn

class EnergyMLP(nn.Module):
    """
        E(x) = <E_theta(x), x> where E_theta is an MLP with nonlinearities.

        x is HVG, log-normalized expression (optionally mean-centered).

        We also expose a latent representation z(x) from the last hidden layer,
        which can be used as a geometry for manifold learning.
    """

    def __init__(self,
                 n_genes: int,
                 hidden_dims: Sequence[int] = (512, 512, 512, 512),
                 activation: Optional[nn.Module] = None,
                 ) :
        super().__init__()
        if activation is None:
            activation = nn.Softplus()

        layers = []
        in_dim = n_genes
        for h in hidden_dims:
            layers.append(nn.Linear(in_dim, h))
            layers.append(activation)
            in_dim = h

        # encoder: maps x \in R^{n_genes} to z \in R^{hidden_dims[-1]}
        self.hidden = nn.Sequential(*layers)

        # head: maps z \in R^{hidden_dims[-1]} to v(z) \in R^{n_genes}
        self.vector_head = nn.Linear(in_dim, n_genes)

        # store for convenience
        self.n_genes = n_genes
        self.latent_dim = in_dim

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """
            Standard forward used in training:
            x: (B, D)
            returns: energy (B, )
        """
        if x.dim() == 1:
            x = x.unsqueeze(0)

        z = self.hidden(x)                      # (B, latent_dim)
        v = self.vector_head(z)                  # (B, D)
        energy = (v * x).sum(dim=-1)           # <v(x), x>
        return energy

    @torch.no_grad()
    def encode(self, x: torch.Tensor) -> torch.Tensor:
        """
            Return latent representation z(x) from the last hidden layer.
            x: (B, D)
            returns: z (B, latent_dim)
        """
        if x.dim() == 1:
            x = x.unsqueeze(0)
        z = self.hidden(x)
```

```
    return z

def score(self, x: torch.Tensor) -> torch.Tensor:
    """
    score(x) = \log p(x) = -E(x)
    """
    x = x.clone().detach().requires_grad_(True)
    energy = self.forward(x) # (B, )
    energy_sum = energy.sum()
    (grad,) = torch.autograd.grad(
        energy_sum,
        x,
        create_graph=False,
        retain_graph=False,
        only_inputs=True,
    )
    score = -grad
    return score
```

```
# src/mantra/eggfm/run_energy.py
from __future__ import annotations

from pathlib import Path
from typing import Dict, Any

import numpy as np
import torch
import scanpy as sc
import yaml

from mantra.eggfm.config import EnergyModelConfig, EnergyTrainConfig
from mantra.eggfm.trainer import train_energy_model
from mantra.utils import subset_anndata

def run_energy_training(
    params_path: Path,
    ad_path: Path,
    out_dir: Path,
    space: str = "hvg",
) -> Path:
    """
        High-level entrypoint: load QC\200\231d AnnData, subset HVGs, train EGGFM, save checkpoint.

    Returns the checkpoint Path.
    """
    out_dir.mkdir(parents=True, exist_ok=True)

    params: Dict[str, Any] = yaml.safe_load(params_path.read_text())
    model_cfg = EnergyModelConfig(**params.get("eggfm_model", {}))
    train_cfg = EnergyTrainConfig(**params.get("eggfm_train", {}))

    # 1) Load prepped K562 AnnData
    ad = sc.read_h5ad(str(ad_path))

    # 2) optional subsample for this experiment
    train_n_cells = params["eggfm_train"].get("n_cells_sample", None)
    if train_n_cells is not None:
        ad_prep = subset_anndata(ad, train_n_cells, random_state=params.get("seed", 0))
    else:
        ad_prep = ad

    # 3) restrict to HVGs if present
    if "highly_variable" in ad_prep.var:
        ad_prep = ad_prep[:, ad_prep.var["highly_variable"]].copy()
        # further clamp HVGs to top N by dispersions_norm
        max_hvg = params["eggfm_train"].get("max_hvg", None)
        if max_hvg is not None and ad_prep.n_vars > max_hvg:
            if "dispersions_norm" in ad_prep.var:
                disp = ad_prep.var["dispersions_norm"].to_numpy()
                order = np.argsort(disp)[::-1] # descending
            else:
                # fallback: arbitrary but deterministic
                order = np.arange(ad_prep.n_vars)

            keep_idx = order[:max_hvg]
            ad_prep = ad_prep[:, keep_idx].copy()
            print(
                f"Subsetting HVGs from {len(order)} \u2192 {ad_prep.n_vars} "
                f"(top {max_hvg} by dispersions_norm)"
            )
        else:
            print("No 'highly_variable' flag in ad.var; using all genes as-is.")

    # 4) Train energy model
    bundle = train_energy_model(
```

```
ad_prep=ad_prep,
model_cfg=model_cfg,
train_cfg=train_cfg,
latent_space=space,
)

if bundle.space != "hvg":
    raise ValueError(
        f"Expected energy model in HVG space, got {bundle.space!r}. "
        "For embedding ablations, use a separate experimental path."
    )

energy_model = bundle.model
mean = bundle.mean
std = bundle.std
var_names = bundle.feature_names

# 5) Save checkpoint
ckpt = {
    "state_dict": energy_model.state_dict(),
    "model_cfg": {
        "hidden_dims": list(model_cfg.hidden_dims),
    },
    "n_genes": energy_model.n_genes,
    "var_names": var_names,
    "mean": mean,
    "std": std,
    "space": bundle.space,
}
space_tag = str(bundle.space).replace("X_", "") # e.g. "hvg", "pca"
n_genes = int(energy_model.n_genes)

ckpt_name = f"eggfm_energy_k562_{space_tag}_hvg{n_genes}.pt"
ckpt_path = out_dir / ckpt_name

torch.save(ckpt, ckpt_path)
print(f"Saved EGGFM energy checkpoint to {ckpt_path}", flush=True)

return ckpt_path
```

```
src/mantra/eggfm/trainer.py           Mon Dec 08 18:21:53 2025      1
# src/mantra/eggfm/trainer.py

from __future__ import annotations

from typing import Optional, Dict

import numpy as np
import torch
import scanpy as sc
from torch import optim
from torch.utils.data import DataLoader

from mantra.eggfm.models import EnergyMLP
from mantra.eggfm.dataset import AnnDataExpressionDataset
from mantra.eggfm.config import EnergyModelConfig, EnergyTrainConfig, EnergyModelBundle


class EnergyTrainer:
    """
    Denoising score-matching trainer for EnergyMLP.

    Given:
    - model
    - standardized dataset (AnnDataExpressionDataset)
    - EnergyTrainConfig

    It runs the DSM loop and returns the best-trained model.
    """

    def __init__(
        self,
        model: EnergyMLP,
        dataset: AnnDataExpressionDataset,
        train_cfg: EnergyTrainConfig,
    ) -> None:
        self.model = model
        self.dataset = dataset
        self.train_cfg = train_cfg

        device_str = train_cfg.device or ("cuda" if torch.cuda.is_available() else "cpu")
        self.device = torch.device(device_str)
        self.model.to(self.device)

        self.loader = DataLoader(
            dataset,
            batch_size=train_cfg.batch_size,
            shuffle=True,
            drop_last=True,
        )

        self.optimizer = optim.Adam(
            self.model.parameters(),
            lr=float(train_cfg.lr),           # force-cast in case YAML gave a string
            weight_decay=float(train_cfg.weight_decay),
        )
        self.best_model_state: Optional[Dict[str, torch.Tensor]] = None
        self.best_trait_state: Optional[Dict[str, torch.Tensor]] = None

        self.best_loss: float = float("inf")
        self.best_state_dict: Optional[dict] = None

    def train(self) -> EnergyMLP:
        sigma = float(self.train_cfg.sigma)
        grad_clip = float(self.train_cfg.grad_clip)
        early_stop_patience = int(self.train_cfg.early_stop_patience)
        early_stop_min_delta = float(self.train_cfg.early_stop_min_delta)
        num_epochs = int(self.train_cfg.num_epochs)
```

```
    self.model.train()
    epochs_without_improve = 0

    n_total = len(self.dataset)

    for epoch in range(num_epochs):
        running_loss = 0.0

        for xb in self.loader:
            xb = xb.to(self.device) # (B, D), already standardized

            # Sample Gaussian noise
            eps = torch.randn_like(xb)
            y = xb + sigma * eps
            y.requires_grad_(True)

            # Energy and score
            energy = self.model(y) # (B, )
            energy_sum = energy.sum() # scalar

            (grad_y,) = torch.autograd.grad(
                energy_sum,
                y,
                create_graph=True,
                retain_graph=True,
                only_inputs=True,
            )
            s_theta = -grad_y

            # DSM target: -(y - x) / sigma^2
            target = -(y - xb) / (sigma**2)

            # MSE over batch and dimensions
            loss = ((s_theta - target) ** 2).sum(dim=1).mean()

            self.optimizer.zero_grad()
            loss.backward()
            if grad_clip > 0.0:
                torch.nn.utils.clip_grad_norm_(
                    self.model.parameters(),
                    grad_clip,
                )
            self.optimizer.step()

            running_loss += loss.item() * xb.size(0)
            epoch_loss = running_loss / n_total
            print(
                f"[Energy DSM] Epoch {epoch+1}/{num_epochs} loss={epoch_loss:.6e}",
                flush=True,
            )

            improved = epoch_loss + early_stop_min_delta < self.best_loss
            if improved:
                self.best_loss = epoch_loss
                self.best_state_dict = self.model.state_dict()
                epochs_without_improve = 0
            else:
                epochs_without_improve += 1

            if early_stop_patience > 0 and epochs_without_improve >= early_stop_patience:
                print(
                    f"[Energy DSM] Early stopping at epoch {epoch+1} "
                    f"(best_loss={self.best_loss:.6e})",
                    flush=True,
                )
                break

        if self.best_state_dict is not None:
```

```
        self.model.load_state_dict(self.best_state_dict)

    return self.model

# -----
# High-level convenience wrapper: AnnData -> EnergyModelBundle
# -----


# src/mantra/eggfm/trainer.py

def train_energy_model(
    ad_prep: sc.AnnData,
    model_cfg: EnergyModelConfig,
    train_cfg: EnergyTrainConfig,
    latent_space: str = "hvg",
) -> EnergyModelBundle:
    """
    AnnData -> AnnDataExpressionDataset -> EnergyMLP -> EnergyTrainer.

    latent_space:
        - "hvg": use ad_prep.X (HVG log-normalized expression)
        - anything else: use ad_prep.obsm[latent_space] (e.g. "X_pca", "X_phate")
    """
    # ----- dataset: HVG or embedding -----
    if latent_space == "hvg":
        X = ad_prep.X
        # true gene names in this space
        feature_names = np.array(ad_prep.var_names.astype(str))
    else:
        if latent_space not in ad_prep.obsm:
            raise KeyError(
                f'Requested latent_space={latent_space!r}, '
                f'but it is not in ad_prep.obsm. '
                f'Available keys: {list(ad_prep.obsm.keys())}'
            )
        X = ad_prep.obsm[latent_space]
    print(f'[EGGFM trainer] Latent_space: {latent_space}', latent_space)

    # synthetic feature names so we *don't* confuse them with genes
    D = X.shape[1]
    feature_names = np.array(
        [f'{latent_space}_{i}' for i in range(D)],
        dtype=str,
    )

dataset = AnnDataExpressionDataset(X)
n_genes = dataset.X.shape[1]

# record normalization (always in the *model feature space*)
mean = dataset.mean # [D]
std = dataset.std # [D]

hidden_dims = tuple(model_cfg.hidden_dims)
model = EnergyMLP(
    n_genes=n_genes,
    hidden_dims=hidden_dims,
)
trainer = EnergyTrainer(
    model=model,
    dataset=dataset,
    train_cfg=train_cfg,
)
best_model = trainer.train()

return EnergyModelBundle(
    model=best_model,
```

```
    mean=mean,  
    std=std,  
    feature_names=feature_names,  
    space=latent_space,  
)
```

```
src/mantra/eggfm//__init__.py           Sun Dec 07 02:07:08 2025      1
```

```
# src/mantra/eggfm/__init__.py
```

```
from .models import EnergyMLP
from .dataset import AnnDataExpressionDataset
from .trainer import EnergyTrainer, train_energy_model
from .inference import EnergyScorer

__all__ = [
    "EnergyMLP",
    "AnnDataExpressionDataset",
    "EnergyTrainer",
    "train_energy_model",
    "EnergyScorer",
]
```